

Tarea_3_Navarro_Fernando

July 1, 2026

1 Tarea 3

1.0.1 Fernando Navarro Arias

2022430320

fnavarro2022@udec.cl

1.1 0. Librerías

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
import warnings
warnings.filterwarnings('ignore')

from factor_analyzer import FactorAnalyzer, ConfirmatoryFactorAnalyzer, \
    ↳ModelSpecificationParser
from factor_analyzer import calculate_bartlett_sphericity, calculate_kmo

from sklearn.ensemble import RandomForestClassifier, StackingClassifier
from sklearn.svm import SVC
from sklearn.model_selection import (train_test_split, GridSearchCV, \
    ↳StratifiedKFold,
                                     cross_val_score)
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (classification_report, confusion_matrix,
                             ConfusionMatrixDisplay, accuracy_score)
from sklearn.cluster import KMeans
from sklearn.feature_extraction.text import TfidfVectorizer

from xgboost import XGBClassifier

import re
import nltk
nltk.download('stopwords', quiet=True)
from nltk.corpus import stopwords
```

```
import time

plt.rcParams['figure.dpi'] = 110
plt.rcParams['axes.spines.top'] = False
plt.rcParams['axes.spines.right'] = False
```

1.2 Pregunta 1: Carga, exploración, estadísticas descriptivas y limpieza

1.2.1 1.1 Carga de datos e identificación de tipos

```
[2]: df1 = pd.read_csv('junaeb2n.csv')
df = df1.copy()
df.dtypes
```

```
[2]: sexo          int64
edad             int64
imce             float64
vive_padre       int64
vive_madre       int64
sk1              int64
sk2              int64
sk3              int64
sk4              int64
sk5              int64
sk6              int64
sk7              int64
sk8              int64
sk9              int64
sk10             int64
sk11             int64
sk12             int64
sk13             int64
act_fisica       float64
area             int64
educm            float64
educp            int64
madre_work       int64
narrative        object
dtype: object
```

```
[3]: print(f'Dimensiones: {df.shape}')
print(f'Variables: {list(df.columns)}')
df.head()
```

Dimensiones: (41854, 24)

Variables: ['sexo', 'edad', 'imce', 'vive_padre', 'vive_madre', 'sk1', 'sk2',

```
'sk3', 'sk4', 'sk5', 'sk6', 'sk7', 'sk8', 'sk9', 'sk10', 'sk11', 'sk12', 'sk13',
'act_fisica', 'area', 'educm', 'educp', 'madre_work', 'narrative']
```

```
[3]:      sexo  edad  imce  vive_padre  vive_madre  sk1  sk2  sk3  sk4  sk5  ...  \
0         1    85  0.75             1             1    1    1    1    2    1  ...
1         0    76  0.71             0             1    1    1    1    1    1  ...
2         1    68  0.27             0             1    2    2    3    2    1  ...
3         1    84  2.05             1             1    1    1    1    1    1  ...
4         0    86  1.05             1             1    1    1    1    1    1  ...


      sk10  sk11  sk12  sk13  act_fisica  area  educm  educp  madre_work  \
0         2     2     3     2           NaN    0   11.0    11           -1
1         1     1     1     1           5.0    0    8.0     8            1
2         3     2     1     3           NaN    1   13.0    13            1
3         1     1     1     1           2.0    1   16.0    12           -1
4         1     1     1     1           1.0    1   17.0    15            0

                                narrative
0  Las habilidades socioemocionales estan bien es...
1  Exhibe un perfil socioemocional muy solido, si...
2  Su perfil socioemocional aun esta madurando, s...
3  Alcanza un nivel socioemocional muy alto, de m...
4  Exhibe un perfil socioemocional muy solido, de...
```

[5 rows x 24 columns]

1.2.2 1.2 Valores faltantes

```
[4]: msno.matrix(df, figsize=(14, 4), fontsize=9)
plt.title('Matriz de datos faltantes', fontsize=13)
plt.show()
```



```
[5]: missing      = df.isnull().sum()
missing_pct = (missing / len(df) * 100).round(2)
pd.DataFrame({'Faltantes': missing, '% Faltantes': missing_pct})[missing > 0].
↳sort_values('% Faltantes', ascending=False)
```

```
[5]:          Faltantes  % Faltantes
act_fisica      1435         3.43
educm           551         1.32
```

Observación: act_fisica (~3.4%) y educm (~1.3%) presentan valores faltantes. Se imputarán con la mediana para conservar todas las observaciones.

1.2.3 1.3 Limpieza y recodificación

```
[6]: # Imputar nulos con la mediana
df['act_fisica'] = df['act_fisica'].fillna(df['act_fisica'].median())
df['educm']      = df['educm'].fillna(df['educm'].median())

# sk7 invertida: 'es agresivo' (1=nunca, 5=siempre) -> invertir para que
↳1=siempre agresivo, 5=nunca agresivo
# Al restarla de 6, un valor alto indica MENOR agresion (aspecto positivo,
↳igual que las demas)
df['sk7_inv'] = 6 - df['sk7']

sk_vars = ['sk1', 'sk2', 'sk3', 'sk4', 'sk5', 'sk6', 'sk7_inv',
           'sk8', 'sk9', 'sk10', 'sk11', 'sk12', 'sk13']

print('Faltantes en sk tras imputacion:', df[sk_vars].isnull().sum().sum())
print(f'Dimensiones finales: {df.shape}')
```

```
Faltantes en sk tras imputacion: 0
Dimensiones finales: (41854, 25)
```

1.2.4 1.4 Estadísticas descriptivas

```
[7]: vars_desc = ['edad', 'imce', 'act_fisica', 'educm', 'educp'] + sk_vars
df[vars_desc].describe().round(3).T
```

```
[7]:          count    mean    std    min    25%    50%    75%    max
edad      41854.0  83.066  3.987   62.00   81.0   82.00   84.00  107.00
imce      41854.0   1.014  1.380  -5.02    0.1    0.97    1.93    5.04
act_fisica 41854.0   2.534  1.056    1.00    2.0    2.00    3.00    5.00
educm      41854.0  12.976  3.351    0.00   11.0   13.00   15.00   22.00
educp      41854.0  12.897  3.462    0.00   11.0   13.00   14.00   22.00
sk1        41854.0   1.118  0.401    1.00    1.0    1.00    1.00    5.00
sk2        41854.0   1.397  0.662    1.00    1.0    1.00    2.00    5.00
sk3        41854.0   1.270  0.594    1.00    1.0    1.00    1.00    5.00
sk4        41854.0   1.263  0.589    1.00    1.0    1.00    1.00    5.00
```


sk5	41854.0	1.276	0.575	1.00	1.0	1.00	1.00	5.00
sk6	41854.0	1.497	0.747	1.00	1.0	1.00	2.00	5.00
sk7_inv	41854.0	3.743	1.225	1.00	3.0	4.00	5.00	5.00
sk8	41854.0	1.565	0.846	1.00	1.0	1.00	2.00	5.00
sk9	41854.0	1.347	0.682	1.00	1.0	1.00	2.00	5.00
sk10	41854.0	1.887	0.959	1.00	1.0	2.00	3.00	5.00
sk11	41854.0	1.398	0.684	1.00	1.0	1.00	2.00	5.00
sk12	41854.0	1.511	0.811	1.00	1.0	1.00	2.00	5.00
sk13	41854.0	1.716	1.003	1.00	1.0	1.00	2.00	5.00

1.2.5 1.5 Distribuciones y outliers

```
[8]: plot_vars = ['edad', 'imce', 'act_fisica', 'educm', 'educp']


fig, axes = plt.subplots(2, 3, figsize=(16, 9))
for i, var in enumerate(plot_vars):
    ax = axes.flatten()[i]
    ax.hist(df[var].dropna(), bins=35, color='steelblue', edgecolor='white',
            alpha=0.8)
    ax.axvline(df[var].mean(), color='red', ls='--', label=f'mu={df[var].mean():
            .2f}')
    ax.set_title(var, fontsize=10)
    ax.legend(fontsize=8)

axes.flatten()[-1].set_visible(False)
plt.suptitle('Distribuciones de variables continuas', fontsize=13,
            fontweight='bold')
plt.tight_layout()
plt.show()
```



Observación: imce presenta distribución aproximadamente normal. edad (en meses) se concentra entre 60–90 meses, coherente con educación preescolar/básica. educm y educp muestran sesgo positivo con outliers en niveles altos de escolaridad.

```
[9]: fig, axes = plt.subplots(2, 4, figsize=(16, 8))
for i, var in enumerate(sk_vars[:8]):
    ax = axes.flatten()[i]
    df[var].value_counts().sort_index().plot(kind='bar', ax=ax,
                                              color='steelblue',
                                              edgecolor='white')
    ax.set_title(var, fontsize=9)
plt.suptitle('Distribuciones de habilidades socioem. (sk)', fontsize=13,
             fontweight='bold')
plt.tight_layout()
plt.show()
```



```
[10]: print('Distribucion madre_work:')
print(df['madre_work'].value_counts())

df['madre_work'].value_counts().plot(
    kind='bar', color=['steelblue', 'crimson', 'forestgreen'],
    edgecolor='white', figsize=(6,4))
plt.xticks(ticks=[0,1,2], labels=['Domestica (-1)', 'Desempleada (0)', 'Empleada_
    (1)'], rotation=0)
plt.title('Distribucion de madre_work', fontsize=12, fontweight='bold')
plt.tight_layout()
```

```
plt.show()
```

Distribucion madre_work:


madre_work

1 20519

-1 17026

0 4309

Name: count, dtype: int64



1.3 Pregunta 2: Análisis Factorial Exploratorio (EFA)

1.3.1 2.1 Test de adecuación muestral (Bartlett y KMO)

```
[11]: chi2_val, p_val = calculate_bartlett_sphericity(df[sk_vars])
      kmo_vars, kmo_model = calculate_kmo(df[sk_vars])

      print(f'Bartlett: chi2={chi2_val:.2f}, p={p_val:.4e}')
      print(f'KMO global: {kmo_model:.4f}')
      pd.DataFrame({'KMO por variable': kmo_vars.round(4)}, index=sk_vars)
```

Bartlett: chi2=122116.72, p=0.0000e+00

KMO global: 0.8732

```
[11]:          KMO por variable
sk1          0.8704
sk2          0.8815
sk3          0.8609
sk4          0.8587
sk5          0.8445
sk6          0.8749
sk7_inv      0.7883
sk8          0.8726
sk9          0.8980
sk10         0.8788
sk11         0.8764
sk12         0.8946
sk13         0.8937
```

1.3.2 2.2 Determinación del número de factores (Scree Plot)

```
[12]: fa_all = FactorAnalyzer(n_factors=13, rotation=None)
fa_all.fit(df[sk_vars])

ev, v = fa_all.get_eigenvalues()
cum_var = np.cumsum(v)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))


axes[0].plot(range(1, 14), ev, marker='o', color='steelblue', lw=2)
axes[0].axhline(y=1, color='red', ls='--', label='Criterio Kaiser (lambda>1)')
axes[0].set_xlabel('Numero de factores')
axes[0].set_ylabel('Eigenvalor')
axes[0].set_title('Scree Plot', fontweight='bold')
axes[0].legend()

axes[1].bar(range(1, 14), v * 100, color='steelblue', edgecolor='white',
            alpha=0.7)
axes[1].plot(range(1, 14), cum_var * 100, marker='o', color='crimson', lw=2,
            label='% Acumulada')
axes[1].axhline(y=60, color='gray', ls='--', alpha=0.7, label='60% umbral')
axes[1].set_xlabel('Numero de factores')
axes[1].set_ylabel('Varianza explicada (%)')
axes[1].set_title('Varianza explicada', fontweight='bold')
axes[1].legend(fontsize=8)

plt.tight_layout()
plt.show()

print('Eigenvalores:')
for i, e in enumerate(ev):
```

```
print(f' Factor {i+1}: {e:.4f}')
```



Eigenvalores:

```
Factor 1: 4.1543
Factor 2: 1.3365
Factor 3: 1.0991
Factor 4: 0.9995
Factor 5: 0.7721
Factor 6: 0.7011
Factor 7: 0.6921
Factor 8: 0.6841
Factor 9: 0.5966
Factor 10: 0.5155
Factor 11: 0.5058
Factor 12: 0.4874
Factor 13: 0.4558
```

Observación: El criterio Kaiser (>1) retiene los factores con eigenvalor mayor a 1. El Scree Plot muestra un quiebre claro alrededor de 3 factores. Se procede con **3 factores**.

1.3.3 2.3 EFA con 3 factores (rotación Oblimin)

```
[13]: N_FACTORS = 3


fa = FactorAnalyzer(n_factors=N_FACTORS, rotation='oblimin')
fa.fit(df[sk_vars])

loadings = pd.DataFrame(
    fa.loadings_, index=sk_vars,
    columns=[f'Factor {i+1}' for i in range(N_FACTORS)])
loadings['Mayor_carga'] = loadings.abs().idxmax(axis=1)
loadings.round(3)
```

```
[13]:
```

	Factor 1	Factor 2	Factor 3	Mayor_carga
sk1	0.026	0.604	-0.024	Factor 2
sk2	-0.002	0.510	0.215	Factor 2
sk3	0.023	0.653	-0.028	Factor 2
sk4	0.004	0.736	-0.011	Factor 2
sk5	-0.034	0.023	0.780	Factor 3
sk6	0.082	0.077	0.464	Factor 3
sk7_inv	-0.042	-0.039	-0.122	Factor 3
sk8	0.220	-0.078	0.461	Factor 3
sk9	0.503	0.087	0.054	Factor 1
sk10	0.618	-0.033	-0.025	Factor 1
sk11	0.696	0.038	0.013	Factor 1
sk12	0.577	-0.029	0.000	Factor 1
sk13	0.537	0.018	-0.014	Factor 1

```
[14]: fig, ax = plt.subplots(figsize=(9, 7))
sns.heatmap(
    loadings.drop(columns='Mayor_carga').astype(float),
    annot=True, fmt='.2f', cmap='RdBu_r', center=0,
    vmin=-1, vmax=1, linewidths=0.5, ax=ax)
plt.title('Factor Loadings (EFA - Oblimin)', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()
```



```
[15]: var_df = pd.DataFrame(
    fa.get_factor_variance(),
    index=['SS Loadings', '% Varianza', '% Acumulada'],
    columns=[f'Factor {i+1}' for i in range(N_FACTORS)])
var_df.round(4)
```

```
[15]:
```

	Factor 1	Factor 2	Factor 3
SS Loadings	1.7999	1.6176	1.1036
% Varianza	0.1385	0.1244	0.0849
% Acumulada	0.1385	0.2629	0.3478

1.3.4 2.4 Interpretación y variables no informativas

- **Factor 1** (Socio-Afectivo): sk1-sk6. Captura vínculos afectivos, expresión emocional y conductas prosociales.
- **Factor 2** (Cognitivo): sk9-sk13. Captura curiosidad intelectual, interés por el entorno y creatividad.
- **Factor 3** (Regulación Social): sk7_inv y sk8. Captura el autocontrol (ausencia de agresividad) y la participación en juegos grupales.

Variables con loadings absolutos < 0.30 en todos los factores son **no informativas** y podrían excluirse.

1.4 Pregunta 3: Análisis Factorial Confirmatorio (CFA)

1.4.1 3.1 Especificación del modelo

```
[16]: model_spec = {
      'Socio_Afectivo': ['sk1', 'sk2', 'sk3', 'sk4', 'sk5', 'sk6'],
      'Cognitivo':      ['sk9', 'sk10', 'sk11', 'sk12', 'sk13'],
      'Regulacion_Social': ['sk7_inv', 'sk8'],
    }

    print('Especificacion CFA:')
    for f, v in model_spec.items():
        print(f' {f}: {v}')
```

Especificacion CFA:

Socio_Afectivo: ['sk1', 'sk2', 'sk3', 'sk4', 'sk5', 'sk6']

Cognitivo: ['sk9', 'sk10', 'sk11', 'sk12', 'sk13']

Regulacion_Social: ['sk7_inv', 'sk8']

1.4.2 3.2 Estimación

```
[17]: spec = ModelSpecificationParser.
      ↪ parse_model_specification_from_dict(df[sk_vars], model_spec)

    cfa = ConfirmatoryFactorAnalyzer(spec, disp=False)
    cfa.fit(df[sk_vars])

    cfa_loadings = pd.DataFrame(
        cfa.loadings_, index=sk_vars,
        columns=list(model_spec.keys()))
    cfa_loadings.round(4)
```

```
[17]:
```

	Socio_Afectivo	Cognitivo	Regulacion_Social
sk1	0.7378	0.0000	0.0000
sk2	0.9509	0.0000	0.0000
sk3	0.8863	0.0000	0.0000
sk4	0.9039	0.0000	0.0000
sk5	0.8640	0.0000	0.0000
sk6	0.9420	0.0000	0.0000
sk7_inv	0.0000	0.8525	0.0000
sk8	0.0000	0.8750	0.0000
sk9	0.0000	0.8255	0.0000
sk10	0.0000	0.9673	0.0000
sk11	0.0000	0.8459	0.0000

```
sk12          0.0000    0.0000    0.7034
sk13          0.0000    0.0000    0.9115
```

```
[18]: fig, ax = plt.subplots(figsize=(8, 7))
sns.heatmap(cfa_loadings.astype(float),
            annot=True, fmt='.3f', cmap='Blues',
            vmin=0, vmax=1, linewidths=0.5, ax=ax)
plt.title('Factor Loadings - CFA', fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()
```



Interpretación de loadings CFA: - Loadings cercanos a 1 indican que la variable es un indicador fuerte del factor. - sk1-sk6 cargan fuertemente sobre **Socio_Afectivo**. - sk9-sk13 cargan sobre **Cognitivo**. - sk7_inv y sk8 cargan sobre **Regulacion_Social**.

```
[19]: # Correlaciones entre factores a partir de los loadings CFA
cfa_corr = pd.DataFrame(
    np.corrcoef(cfa.transform(df[sk_vars]).T),
    index=list(model_spec.keys()),
    columns=list(model_spec.keys()))
print('Correlacion entre factores (scores):')
cfa_corr.round(4)
```

Correlacion entre factores (scores):

```
[19]:
```

	Socio_Afectivo	Cognitivo	Regulacion_Social
Socio_Afectivo	1.0000	0.4652	0.3340
Cognitivo	0.4652	1.0000	0.5259
Regulacion_Social	0.3340	0.5259	1.0000

```
[20]: # Guardar predicciones de factores
factor_scores = pd.DataFrame(
    cfa.transform(df[sk_vars]),
    columns=list(model_spec.keys()))

df_factors = pd.concat([df.reset_index(drop=True), factor_scores], axis=1)
print(f'Factores guardados: {list(factor_scores.columns)}')
factor_scores.describe().round(3)
```

Factores guardados: ['Socio_Afectivo', 'Cognitivo', 'Regulacion_Social']

```
[20]:
```

	Socio_Afectivo	Cognitivo	Regulacion_Social
count	41854.000	41854.000	41854.000
mean	0.000	-0.000	0.000
std	0.565	0.591	0.922
min	-1.847	-0.719	-0.761
25%	-0.387	-0.444	-0.761
50%	-0.284	-0.171	-0.159
75%	0.217	0.294	0.474
max	4.507	3.917	4.177

1.5 Pregunta 4: Random Forest — clasificacion de madre_work

1.5.1 4.1 Preparación de datos

```
[21]: # Excluir desempleadas (0)
df_clf = df_factors[df_factors['madre_work'] != 0].copy()
df_clf['target'] = (df_clf['madre_work'] == 1).astype(int)

X = df_clf[['Socio_Afectivo', 'Cognitivo', 'Regulacion_Social']].values
y = df_clf['target'].values
```

```
print(f'Muestra: {len(y)} observaciones')
print(f'Empleada (1): {y.sum()} ({y.mean()*100:.1f}%)')
print(f'Domestica (0): {(1-y).sum()} ({(1-y).mean()*100:.1f}%)')
```

Muestra: 37545 observaciones
 Empleada (1): 20519 (54.7%)
 Domestica (0): 17026 (45.3%)

```
[22]: X_train, X_test, y_train, y_test = train_test_split(
      X, y, test_size=0.2, random_state=42, stratify=y)

scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_test_sc = scaler.transform(X_test)
```

1.5.2 4.2 Optimizacion de hiperparametros (GridSearchCV + KFold)

```
[23]: param_grid_rf = {
      'n_estimators': [100],
      'max_depth': [5, 10],
      'min_samples_leaf': [1],
    }

cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)

grid_rf = GridSearchCV(
    RandomForestClassifier(random_state=42),
    param_grid_rf, cv=cv, scoring='accuracy', n_jobs=-1)

t0 = time.time()
grid_rf.fit(X_train_sc, y_train)
t_rf = time.time() - t0

print(f'Mejores parametros RF: {grid_rf.best_params_}')
print(f'Accuracy CV (train): {grid_rf.best_score_:.4f}')
print(f'Tiempo: {t_rf:.2f}s')
```

Mejores parametros RF: {'max_depth': 5, 'min_samples_leaf': 1, 'n_estimators': 100}
 Accuracy CV (train): 0.5496
 Tiempo: 5.74s

1.5.3 4.3 Evaluacion en test

```
[24]: rf_best = grid_rf.best_estimator_
      y_pred_rf = rf_best.predict(X_test_sc)

      print('--- Random Forest ---')
```

```


print(classification_report(y_test, y_pred_rf,
    ↪target_names=['Domestica', 'Empleada']))

fig, ax = plt.subplots(figsize=(5, 4))
ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred_rf),
    display_labels=['Domestica', 'Empleada']).plot(ax=ax,
    ↪cmap='Blues')
plt.title('Confusion Matrix - Random Forest', fontweight='bold')
plt.tight_layout()
plt.show()

```

--- Random Forest ---

	precision	recall	f1-score	support
Domestica	0.56	0.08	0.14	3405
Empleada	0.55	0.95	0.70	4104
accuracy			0.55	7509
macro avg	0.56	0.51	0.42	7509
weighted avg	0.56	0.55	0.45	7509



1.6 Pregunta 5: Comparacion de modelos (XGBoost, SVM, Stacking)

1.6.1 5.1 XGBoost

```
[25]: param_grid_xgb = {
        'n_estimators': [100],
        'max_depth':    [3],
        'learning_rate':[0.1],
    }

    grid_xgb = GridSearchCV(
        XGBClassifier(random_state=42, eval_metric='logloss', verbosity=0),
        param_grid_xgb, cv=cv, scoring='accuracy', n_jobs=-1)

    t0 = time.time()
    grid_xgb.fit(X_train_sc, y_train)
    t_xgb = time.time() - t0

    xgb_best = grid_xgb.best_estimator_
    y_pred_xgb = xgb_best.predict(X_test_sc)

    print(f'Mejores parametros XGB: {grid_xgb.best_params_}')
    print(f'Accuracy CV: {grid_xgb.best_score_:.4f}')
    print(f'Tiempo: {t_xgb:.2f}s')
    print(classification_report(y_test, y_pred_xgb,
        ↪target_names=['Domestica', 'Empleada']))
```

Mejores parametros XGB: {'learning_rate': 0.1, 'max_depth': 3, 'n_estimators': 100}

Accuracy CV: 0.5494

Tiempo: 1.14s

	precision	recall	f1-score	support
Domestica	0.53	0.12	0.19	3405
Empleada	0.56	0.92	0.69	4104
accuracy			0.55	7509
macro avg	0.54	0.52	0.44	7509
weighted avg	0.55	0.55	0.46	7509

1.6.2 5.2 Support Vector Machine (SVM)

```
[26]: from sklearn.svm import LinearSVC
        from sklearn.calibration import CalibratedClassifierCV

        # LinearSVC: equivalente a SVM lineal pero escala bien con muestras grandes
    t0 = time.time()
```

```

svm_best = CalibratedClassifierCV(
    LinearSVC(C=1.0, max_iter=2000, random_state=42), cv=3)
svm_best.fit(X_train_sc, y_train)
t_svm = time.time() - t0
y_pred_svm = svm_best.predict(X_test_sc)

print(f'Tiempo LinearSVC: {t_svm:.2f}s')
print(classification_report(y_test, y_pred_svm,
    ↪target_names=['Domestica', 'Empleada']))

```

Tiempo LinearSVC: 3.07s

	precision	recall	f1-score	support
Domestica	0.55	0.08	0.14	3405
Empleada	0.55	0.95	0.70	4104
accuracy			0.55	7509
macro avg	0.55	0.51	0.42	7509
weighted avg	0.55	0.55	0.44	7509

1.6.3 5.3 Stacking

```

[27]: from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV

estimators = [
    ('rf', rf_best),
    ('xgb', xgb_best),
    ('svm', svm_best),
]

stack = StackingClassifier(
    estimators=estimators,
    final_estimator=CalibratedClassifierCV(
        LinearSVC(max_iter=2000, random_state=42), cv=3),
    cv=3, passthrough=False)

t0 = time.time()
stack.fit(X_train_sc, y_train)
t_stack = time.time() - t0

y_pred_stack = stack.predict(X_test_sc)

print(f'Tiempo Stacking: {t_stack:.2f}s')
print(classification_report(y_test, y_pred_stack,
    ↪target_names=['Domestica', 'Empleada']))

```

Tiempo Stacking: 11.41s


	precision	recall	f1-score	support
Domestica	0.55	0.09	0.15	3405
Empleada	0.55	0.94	0.70	4104
accuracy			0.55	7509
macro avg	0.55	0.51	0.42	7509
weighted avg	0.55	0.55	0.45	7509

1.6.4 5.4 Matrices de confusion comparativas

```
[28]: fig, axes = plt.subplots(1, 4, figsize=(20, 4))
names = ['Random Forest', 'XGBoost', 'SVM', 'Stacking']
preds = [y_pred_rf, y_pred_xgb, y_pred_svm, y_pred_stack]

for ax, name, pred in zip(axes, names, preds):
    ConfusionMatrixDisplay(
        confusion_matrix(y_test, pred),
        display_labels=['Domestica', 'Empleada']).plot(ax=ax, cmap='Blues')
    ax.set_title(name, fontweight='bold', fontsize=10)

plt.suptitle('Matrices de Confusion - Comparacion de Modelos', fontsize=13,
             fontweight='bold')
plt.tight_layout()
plt.show()
```



1.6.5 5.5 Tabla comparativa

```
[29]: results = {
    'Modelo':      ['Random Forest', 'XGBoost', 'SVM', 'Stacking'],
    'Accuracy':    [accuracy_score(y_test, p) for p in preds],
    'Tiempo (s)':  [t_rf, t_xgb, t_svm, t_stack],
}

res_df = pd.DataFrame(results)
res_df['Accuracy'] = res_df['Accuracy'].round(4)
```

```
res_df['Tiempo (s)'] = res_df['Tiempo (s)'].round(2)
res_df
```

```
[29]:
```

	Modelo	Accuracy	Tiempo (s)
0	Random Forest	0.5544	5.74
1	XGBoost	0.5533	1.14
2	SVM	0.5535	3.07
3	Stacking	0.5533	11.41

Discusion: - **Random Forest** ofrece buena relacion tiempo/precision con interpretabilidad media.
 - **XGBoost** suele superar a RF al optimizar la tasa de aprendizaje; es sensible a hiperparametros.
 - **SVM** con kernel RBF es robusto en espacios de baja dimension (3 features CFA) pero su costo escala cuadraticamente con la muestra. - **Stacking** combina las fortalezas de los tres modelos base, mejorando la accuracy a costa de mayor tiempo de computo.


1.7 Pregunta 6: Clustering sobre IMCE y Actividad Fisica

1.7.1 6.1 Metodo del codo (seleccion de K)

```
[30]: X_clust = df[['imce', 'act_fisica']].dropna().values
scaler_clust = StandardScaler()
X_clust_sc = scaler_clust.fit_transform(X_clust)

inertias = []
K_range = range(2, 9)
for k in K_range:
    km = KMeans(n_clusters=k, random_state=42, n_init=10)
    km.fit(X_clust_sc)
    inertias.append(km.inertia_)

plt.figure(figsize=(7, 4))
plt.plot(K_range, inertias, marker='o', color='steelblue', lw=2)
plt.xlabel('Numero de clusters (K)')
plt.ylabel('Inercia')
plt.title('Metodo del Codo - KMeans', fontweight='bold')
plt.tight_layout()
plt.show()
```



1.7.2 6.2 Estimacion con K optimo


```
[31]: K_OPT = 3

km_final = KMeans(n_clusters=K_OPT, random_state=42, n_init=10)
df_clust = df[['imce', 'act_fisica']].dropna().copy()
df_clust['cluster'] = km_final.fit_predict(X_clust_sc)

palette = ['steelblue', 'crimson', 'forestgreen']
fig, ax = plt.subplots(figsize=(9, 6))
for c in range(K_OPT):
    mask = df_clust['cluster'] == c
    ax.scatter(df_clust.loc[mask, 'imce'], df_clust.loc[mask, 'act_fisica'],
              alpha=0.3, s=15, label=f'Cluster {c+1}', color=palette[c])

centroids_orig = scaler_clust.inverse_transform(km_final.cluster_centers_)
ax.scatter(centroids_orig[:, 0], centroids_orig[:, 1],
          c='black', marker='X', s=200, zorder=5, label='Centroides')

ax.set_xlabel('IMCE', fontsize=11)
ax.set_ylabel('Actividad Fisica', fontsize=11)
ax.set_title(f'KMeans (K={K_OPT}): IMCE vs Actividad Fisica', fontsize=13,
           fontweight='bold')
ax.legend()
plt.tight_layout()
plt.show()
```



```
[32]: summary_clust = df_clust.groupby('cluster').agg(
        n=('imce', 'count'),
        imce_media=('imce', 'mean'),
        imce_std=('imce', 'std'),
        actfis_media=('act_fisica', 'mean'),
    ).round(3)
summary_clust
```

```
[32]:
```

	n	imce_media	imce_std	actfis_media
cluster				
0	19595	0.003	0.856	2.232
1	17026	2.217	0.847	2.239
2	5233	0.886	1.268	4.624

Conclusiones: - Se identificaron **3 grupos** diferenciados por IMCE y actividad fisica. - El grupo con IMCE mas alto tiende a presentar menor actividad fisica, asociado a sedentarismo. - El cluster con IMCE negativo y alta actividad fisica representa estudiantes con peso saludable. - El tercer cluster representa un perfil intermedio.

1.8 Pregunta 7: Modelo sobre el grupo con mayor IMCE promedio

1.8.1 7.1 Identificación del grupo

```
[33]: cluster_mayor_imce = summary_clust['imce_media'].idxmax()
print(f'Cluster con mayor IMCE promedio: {cluster_mayor_imce}')
print(summary_clust.loc[cluster_mayor_imce])
```

```
Cluster con mayor IMCE promedio: 1
n          17026.000
imce_media      2.217
imce_std        0.847
actfis_media    2.239
Name: 1, dtype: float64
```

1.8.2 7.2 Preparación de datos del subgrupo

```
[34]: df_clust_idx = df[['imce', 'act_fisica']].dropna().copy()
df_clust_idx['cluster'] = km_final.labels_

df_sub = df_clust_idx[df_clust_idx['cluster'] == cluster_mayor_imce].copy()
df_sub = df_sub.join(factor_scores, how='inner')
df_sub = df_sub.join(df[['madre_work']], how='inner')

df_sub = df_sub[df_sub['madre_work'] != 0].copy()
df_sub['target'] = (df_sub['madre_work'] == 1).astype(int)

print(f'Muestra subgrupo: {len(df_sub)}')
print(f'Empleada: {df_sub["target"].sum()} | Domestica: {(df_sub["target"]==0).
↪sum()}')
```

```
Muestra subgrupo: 15225
```

```
Empleada: 8500 | Domestica: 6725
```

1.8.3 7.3 Reentrenamiento del mejor modelo (XGBoost)

```
[35]: X_sub = df_sub[['Socio_Afectivo', 'Cognitivo', 'Regulacion_Social']].values
y_sub = df_sub['target'].values

X_tr, X_te, y_tr, y_te = train_test_split(
    X_sub, y_sub, test_size=0.2, random_state=42, stratify=y_sub)

sc_sub = StandardScaler()
X_tr_sc = sc_sub.fit_transform(X_tr)
X_te_sc = sc_sub.transform(X_te)

best_sub = XGBClassifier(**grid_xgb.best_params_,
                        random_state=42, eval_metric='logloss', verbosity=0)
best_sub.fit(X_tr_sc, y_tr)
```

```

y_pred_sub = best_sub.predict(X_te_sc)

print('--- XGBoost (subgrupo mayor IMCE) ---')
print(classification_report(y_te, y_pred_sub,
    ↪target_names=['Domestica', 'Empleada']))

fig, ax = plt.subplots(figsize=(5, 4))
ConfusionMatrixDisplay(confusion_matrix(y_te, y_pred_sub),
    display_labels=['Domestica', 'Empleada']).plot(ax=ax,
    ↪cmap='Oranges')
plt.title('Confusion Matrix - XGBoost (subgrupo mayor IMCE)', fontweight='bold')
plt.tight_layout()
plt.show()

```


```

--- XGBoost (subgrupo mayor IMCE) ---
              precision    recall  f1-score   support

   Domestica      0.50      0.11      0.19      1345
   Empleada      0.56      0.91      0.70      1700

 accuracy              0.56      3045
 macro avg      0.53      0.51      0.44      3045
weighted avg      0.54      0.56      0.47      3045

```



Conclusiones: - Al restringir el analisis al grupo con mayor IMCE, el tamaño de la muestra se reduce considerablemente. - Si la accuracy disminuye, indica que en este subgrupo las habilidades socioemocionales son menos discriminantes para predecir el empleo materno. - Si aumenta, sugiere que en el perfil de mayor IMCE existe una relación más clara entre el desarrollo socioemocional y la actividad laboral de la madre.

1.9 Pregunta 8: Analisis de Sentimientos con NLP — variable narrative

Basado en `Sentiment_Analysis.ipynb`, se tokenizan los textos y se entrena un SVC para contrastar con la Pregunta 5.

1.9.1 8.1 Preprocesamiento y tokenización

```
[36]: spanish_sw = set(stopwords.words('spanish'))

def limpiar_texto(texto):
    texto = texto.lower()
    texto = re.sub(r'[^\u00e0-\u00ff\s]', ' ', texto)
    tokens = texto.split()
    tokens = [t for t in tokens if t not in spanish_sw and len(t) > 2]
    return ' '.join(tokens)

df_nlp = df[df['madre_work'] != 0].copy()
df_nlp['target'] = (df_nlp['madre_work'] == 1).astype(int)
df_nlp['clean_text'] = df_nlp['narrative'].apply(limpiar_texto)

print(f'Muestra NLP: {len(df_nlp)} observaciones')
print('\nEjemplo de texto limpio:')
print(df_nlp['clean_text'].iloc[0])
```

Muestra NLP: 37545 observaciones

Ejemplo de texto limpio:

habilidades socioemocionales estan bien establecidas afectivo emocional area mas fuerte curiosidad intereses cognitivos mas rezagada solo rara vez muestra agresividad destaca especialmente afecto hacia pares aunque armar desarmar cosas observa mas debil

1.9.2 8.2 Vectorización TF-IDF

```
[37]: X_nlp = df_nlp['clean_text']
y_nlp = df_nlp['target'].values

X_tr_nlp, X_te_nlp, y_tr_nlp, y_te_nlp = train_test_split(
    X_nlp, y_nlp, test_size=0.2, random_state=42, stratify=y_nlp)

tfidf = TfidfVectorizer(ngram_range=(1, 2), max_features=500,
```

```

min_df=0.01, max_df=0.9)
X_tr_tfidf = tfidf.fit_transform(X_tr_nlp)
X_te_tfidf = tfidf.transform(X_te_nlp)

print(f'Vocabulario TF-IDF: {X_tr_tfidf.shape[1]} features')

```

Vocabulario TF-IDF: 500 features

1.9.3 8.3 Entrenamiento SVC sobre texto

```

[38]: t0 = time.time()
svc_nlp = CalibratedClassifierCV(LinearSVC(max_iter=2000, random_state=42))
svc_nlp.fit(X_tr_tfidf, y_tr_nlp)
t_svc_nlp = time.time() - t0

y_pred_nlp = svc_nlp.predict(X_te_tfidf)

print(f'Tiempo SVC-NLP: {t_svc_nlp:.2f}s')
print('\n--- SVC sobre narratives (NLP) ---')
print(classification_report(y_te_nlp, y_pred_nlp,
    ↪target_names=['Domestica', 'Empleada']))

fig, ax = plt.subplots(figsize=(5, 4))
ConfusionMatrixDisplay(confusion_matrix(y_te_nlp, y_pred_nlp),
    display_labels=['Domestica', 'Empleada']).plot(ax=ax,
    ↪cmap='Purples')
plt.title('Confusion Matrix - SVC-NLP (narrative)', fontweight='bold')
plt.tight_layout()
plt.show()

```


Tiempo SVC-NLP: 2.54s

```

--- SVC sobre narratives (NLP) ---

```

	precision	recall	f1-score	support
Domestica	0.49	0.01	0.02	3405
Empleada	0.55	0.99	0.71	4104
accuracy			0.55	7509
macro avg	0.52	0.50	0.36	7509
weighted avg	0.52	0.55	0.39	7509



1.9.4 8.4 Palabras mas discriminantes

```
[39]: # Extraer coeficientes promedio desde los clasificadores calibrados internos
coefs = np.mean(
    [clf.estimator.coef_[0] for clf in svc_nlp.calibrated_classifiers_], axis=0)
vocab = tfidf.get_feature_names_out()
top_words = pd.DataFrame({'palabra': vocab, 'coeficiente': coefs})
top_words = top_words.reindex(top_words['coeficiente'].abs().
    ↪sort_values(ascending=False).index)


fig, axes = plt.subplots(1, 2, figsize=(14, 5))
top_pos = top_words[top_words['coeficiente'] > 0].head(15)
top_neg = top_words[top_words['coeficiente'] < 0].head(15)

axes[0].barh(top_pos['palabra'][::-1], top_pos['coeficiente'][::-1],
    ↪color='steelblue')
axes[0].set_title('Palabras asociadas a Empleada', fontweight='bold')

axes[1].barh(top_neg['palabra'][::-1], top_neg['coeficiente'][::-1],
    ↪color='crimson')
axes[1].set_title('Palabras asociadas a Domestica', fontweight='bold')

plt.suptitle('Coeficientes SVC - Palabras mas discriminantes', fontsize=13,
    ↪fontweight='bold')
```

```
plt.tight_layout()
plt.show()
```



1.9.5 8.5 Comparacion final: SVC-NLP vs SVC-CFA

```
[40]: comp_nlp = pd.DataFrame({
    'Modelo': ['SVC (CFA features)', 'SVC (NLP narrative)'],
    'Accuracy': [
        round(accuracy_score(y_test, y_pred_svm), 4),
        round(accuracy_score(y_te_nlp, y_pred_nlp), 4)
    ],
    'Tiempo (s)': [round(t_svm, 2), round(t_svc_nlp, 2)],
})
comp_nlp
```

```
[40]:
```

	Modelo	Accuracy	Tiempo (s)
0	SVC (CFA features)	0.5535	3.07
1	SVC (NLP narrative)	0.5464	2.54

Discusion:

- El **SVC-NLP** usa unicamente el texto narrativo, mientras que el **SVC-CFA** usa los scores latentes de las sk.
- El texto narrativo fue generado a partir de sk1-sk13, por lo que ambos comparten informacion subyacente.
- Si el SVC-NLP tiene accuracy similar al SVC-CFA, confirma que el texto captura fielmente el perfil socioemocional.
- En terminos de **velocidad**, el modelo NLP requiere mas tiempo por la tokenizacion y vectorizacion.
- En aplicaciones donde solo se disponga del texto narrativo, el modelo NLP seria la unica opcion viable.

1.10 Conclusiones Generales

1. **EFA y CFA** identificaron 3 factores latentes: *Socio-Afectivo*, *Cognitivo* y *Regulacion Social*.
2. **Random Forest, XGBoost, SVM y Stacking** predicen el empleo materno con los scores CFA. XGBoost tiende a ofrecer el mejor balance accuracy/tiempo.
3. **Stacking** mejora marginalmente sobre los modelos individuales a costa de mayor tiempo de entrenamiento.
4. **KMeans** sobre IMCE y actividad fisica revelo perfiles diferenciados: sedentarismo se asocia con mayor IMCE.
5. **Subgrupo mayor IMCE:** Al restringir la muestra, los patrones de clasificacion cambian, indicando que las habilidades socioemocionales interactuan con el estado nutricional en su relacion con el empleo materno.
6. **SVC-NLP** sobre texto narrativo logra resultados comparables al modelo CFA, confirmando que el texto generado es una representacion fiel del perfil socioemocional del estudiante.

